

MNIST training

```
In [1]: import numpy as onp
import jax.numpy as jnp
import losses

from training import train
from custom_activations import smooth_leaky_relu
#from tb_logger import tensorboard_decorator
from loggers import *
from utils import *
from datasets.dataloaders import load_dataset
```

Setting hyperparameters and loading the dataset

```
In [2]: # hyperparameters
seed      = 42
num_layers = 2

batch_size = 10
epochs     = 25

lr         = 1e-4

bias       = True
trick      = True
#generate  = True

dataset    = "MNIST"
log_dir    = "mnist_run/"

alpha      = 0.01
activation = lambda x: smooth_leaky_relu(x, alpha=alpha)
loss_pdf   = losses.log_pdf_normal

# Load dataset
train_data, val_data, test_data, transform = load_dataset(dataset)
```

Model definition

A model is specified by its `params` and 3 `g` functions. All the 3 functions perform a forward pass through the model, but with some differences:

- `g_dummy` takes 3 arguments: `params`, `dummy_params`, `data`. The `dummy_params` are there only to be able to accumulate the `delta` gradients when needed (refer to Section D of the Supplementary Material).
- `g_layerwise` returns the output `z` (the predicted sources) and the activations `ys` for each layer (no non-linearities applied).
- `g`: returns `z` only

The model constructors are found in `models.py`.

```
In [3]: import models

d = train_data.shape[-1] # data dimensionality

# model definition #####
#####
params, g_dummy, g_layerwise, g = models.g_with_dummies(d, num_layers
,
                                activation = activation,
                                seed = seed,
                                batch_size = batch_size,
                                bias = bias)

/home/fissore/.local/lib/python3.8/site-packages/jax/lib/xla_bridge.p
y:123: UserWarning: No GPU/TPU found, falling back to CPU.
  warnings.warn('No GPU/TPU found, falling back to CPU.')
```

Loss function and gradient computation

```

In [4]: # The loss function requires the `g_dummy` model in order to be able
        # to compute
        # the gradients of the `delta` terms (Section D of the Supplementary
        # Material)

        loss = losses.get_loss_deltas(g_dummy, loss_pdf, activation = activation)
        gradient = losses.natural_grad_deltas(loss, activation, bias=bias, trick=trick)

        # When we use biases, we need to project the gradient on the appropriate submanifold
        # (refer to Section F of the Supplementary Material)

        if bias:
            gradient = gradient_padding(gradient)

        # Logging loss values and execution time #####
        #####
        log_losses_ = losses.losses_logger(loss, loss_pdf, val_data, activation, g_layerwise)

        loss_vs_time = []

        def log_loss(params, epoch):
            _, v = log_losses_(params, epoch)
            l = v[3]
            loss_vs_time.append(l)
            return "Loss", [l]

        log_time = timer()
        log_params, get_params = params_getter(params)

        loggers_list = [log_loss, log_time]

```

Training

```
In [6]: train(params, train_data, gradient,
           epochs = epochs,
           lr = lr,
           batch_size = batch_size,
           loggers = loggers_list,
           log_every = 1
        )
```

```
Epoch 0 [Loss [1404.8104248046875]] [Time [60.657156933040824]]
Epoch 1 [Loss [1381.4111328125]] [Time [58.29753472498851]]
Epoch 2 [Loss [1374.9918212890625]] [Time [58.91373840399319]]
Epoch 3 [Loss [1372.2982177734375]] [Time [59.64869629900204]]
Epoch 4 [Loss [1371.12646484375]] [Time [59.480004040000495]]
Epoch 5 [Loss [1370.330810546875]] [Time [58.69616854301421]]
Epoch 6 [Loss [1369.85205078125]] [Time [58.221025504986756]]
Epoch 7 [Loss [1369.617431640625]] [Time [58.67217283899663]]
Epoch 8 [Loss [1368.893798828125]] [Time [59.51819232001435]]
Epoch 9 [Loss [1368.841796875]] [Time [58.65609293297166]]
Epoch 10 [Loss [1368.6226806640625]] [Time [58.873983016994316]]
Epoch 11 [Loss [1368.42529296875]] [Time [59.738252012990415]]
Epoch 12 [Loss [1368.574951171875]] [Time [59.52137082599802]]
Epoch 13 [Loss [1368.4102783203125]] [Time [56.67055111605441]]
Epoch 14 [Loss [1368.41943359375]] [Time [57.885606968950015]]
Epoch 15 [Loss [1368.22412109375]] [Time [59.27402735204669]]
Epoch 16 [Loss [1368.2725830078125]] [Time [59.60896977398079]]
Epoch 17 [Loss [1368.701171875]] [Time [58.50166714400984]]
Epoch 18 [Loss [1369.190673828125]] [Time [59.28502304799622]]
Epoch 19 [Loss [1368.918212890625]] [Time [58.40523395501077]]
Epoch 20 [Loss [1368.4415283203125]] [Time [58.46523604996037]]
Epoch 21 [Loss [1367.9398193359375]] [Time [59.3193076979951]]
Epoch 22 [Loss [1368.39990234375]] [Time [58.36421982903266]]
Epoch 23 [Loss [1370.017578125]] [Time [57.7569236769923]]
Epoch 24 [Loss [1369.080078125]] [Time [59.772814610973]]
Epoch 25 [Loss [1368.4273681640625]] [Time [58.79686041100649]]
```

```
Out[6]: 0
```

```
In [31]: from matplotlib import pyplot as plt

fig, ax = plt.subplots()
ax.set_ylabel('-log P(x)', fontsize='xx-large')
ax.set_xlabel('Epochs', fontsize='xx-large')
plt.plot(loss_vs_time)
plt.show()
```

